

⌵ Lab 08: Vision Transformers and VLMs, Talking to Your Images

ITAI 1378 | Module 08 Companion Lab | 2026

Author: Patricia McManus **Collaborator:** Claude (Anthropic)

Course: AI and Robotics Program, Digital and Information Technology Center of Excellence, Houston City College

In Module 08 you learned how Vision Transformers chop an image into patches and treat them like words, and how Vision Language Models (VLMs) connect a model's eyes to a language model so it can describe, answer, and reason about images. The frontier VLMs you studied (GPT-5.5, Claude Opus 4.7, the Gemini 3 family) cost money per image. Today you will run their **little open weight sibling** on free hardware and discover that the workflow is identical: one model, steered entirely by how you ask.

By the end of this lab you will be able to:

- 1. Visualize how a Vision Transformer "reads" an image as patches
- 2. Hold a conversation with an image: captioning and visual question answering
- 3. Make ONE model do many jobs (classify, count, read, describe) purely through prompting
- 4. Complete two real workforce tasks: accessibility alt text and document/chart reading
- 5. Audit a VLM's reliability and catch it hallucinating
- 6. Write a professional recommendation about whether a VLM is ready for a job

How This Lab Works

- Cells marked **RUN AS IS** are complete. Run them and read the comments.
- Cells marked **YOUR CODE HERE** have a TODO for you to complete. Tips are in the comments.
- Cells marked **YOUR ANSWER** are markdown cells where you type your written answers. Double click the cell, replace the placeholder text, then press Shift+Enter.
- **YOUR TURN** sections are graded checkpoints. **Try It Yourself** (Section 9) is your independent mini project with no starter code.

Your GPU Budget (read this!)

This lab is even lighter than Lab 07:

- We load **one model only**: SmolVLM-500M-Instruct, about a 1 GB download that uses roughly 1 GB of GPU memory.
- Each question you ask takes a few seconds. The whole lab needs a few dozen questions.
- The CPU fallback is genuinely usable this time. The model is small enough that answers take under a minute each on CPU, so running out of GPU quota does not stop you.
- Plan to finish in one sitting so the downloaded model stays cached.

Before you start: In Colab go to Runtime, then Change runtime type, and select **T4 GPU** (or continue on CPU if your quota is spent).

Grading Rubric (100 points, bonus available)

Implementation, 60 points (six items, 10 points each)

Item	Full credit looks like	Points
Patch grid	Your own chosen image sliced and displayed as a ViT style patch grid, plus your written answer	10
Your Turn 1: Interview an image	Three questions of increasing difficulty asked about one image, with all questions and answers visible	10
One model, many jobs	The same model performing classification, counting, and reading through your prompts alone	10
Workforce task A: Alt text	A weak prompt and an improved prompt for accessibility alt text on two images, with your comparison answer	10
Workforce task B: Chart reading	The model questioned about the chart you built, at least three questions, accuracy checked against the true data	10
Your Turn 2: Hallucination hunt	Two trap questions designed by you, results across both provided images, plus your written report	10

Analysis, 40 points

Item	Full credit looks like
Try It Yourself audit	A chosen workplace scenario (5), a designed prompt tested for reliability across three phrasings (5), a professional deploy or do not deploy memo naming evidence
Reflection questions	All ten answered in complete sentences in your own words, 2 points each

Bonus, up to 20 points: see Section 11.

Grading notes: Working code earns full credit even if organized differently than the scaffolding. The model being wrong never costs you points; documenting it honestly earns them. What is graded is your process and your judgment. There is no deduction for late submission.

Academic integrity: Written answers must be your own words. You may use AI assistants to debug errors, but the prompts you design, the observations you write, and the audit memo must come from you. Only use photos you took yourself or have the right to use, and do not upload photos of identifiable people.

Submission: This notebook is the only deliverable. Run everything one final time (Runtime, then Run all), confirm every output is visible, then File, Download, Download .ipynb **with all outputs visible**. Name it **L08_LastName_FirstName_ITAI1378.ipynb** and submit to Canvas.

Section 1: Setup

RUN AS IS. We install the Hugging Face libraries and detect your hardware. Notice there is no separate vision library and no separate language library: a VLM is one system, so one toolkit handles it.

```
1 # RUN AS IS
2 !pip install -q transformers accelerate
3
4 import torch
5
6 if torch.cuda.is_available():
7     DEVICE = "cuda"
8     DTYPE = torch.float16
9     print("GPU ready:", torch.cuda.get_device_name(0))
10 else:
11     DEVICE = "cpu"
12     DTYPE = torch.float32
13     print("CPU fallback active. This model is small enough that CPU works,")
14     print("just expect each answer to take up to a minute. Be patient!")
```

GPU ready: Tesla T4


Section 2: Seeing Like a Vision Transformer (Zero GPU Cost)

Before we talk to a VLM, let's see what its eyes actually receive. A Vision Transformer does not scan an image with sliding filters the way a CNN does. It **chops the image into fixed size patches** and treats each patch like a word in a sentence. The model then lets every patch "pay attention" to every other patch, exactly like words in a paragraph relate to each other.

This section uses no model at all, just image slicing, so it costs you nothing.

```
1 # RUN AS IS
2 # Download a sample image and slice it into a ViT style patch grid.
3 import urllib.request
4 from PIL import Image
5 import matplotlib.pyplot as plt
6
7 url = "https://huggingface.co/datasets/huggingface/documentation-images/resolve/main/diffusers/cat.png"
8 urllib.request.urlretrieve(url, "sample_cat.png")
9 img = Image.open("sample_cat.png").convert("RGB").resize((384, 384))
10
11 def show_patches(image, patch_size=48):
12     """Slice an image into square patches and display them as a spaced grid,
13     the way a Vision Transformer sees its input."""
14     w, h = image.size
15     cols, rows = w // patch_size, h // patch_size
16     fig, axes = plt.subplots(rows, cols, figsize=(6, 6))
17     for r in range(rows):
18         for c in range(cols):
19             patch = image.crop((c * patch_size, r * patch_size,
20                               (c + 1) * patch_size, (r + 1) * patch_size))
21             axes[r, c].imshow(patch)
22             axes[r, c].axis("off")
23     fig.suptitle(f"{rows * cols} patches, each treated like a word", fontsize=11)
24     plt.tight_layout()
25     plt.show()
26
27 show_patches(img)
```

64 patches, each treated like a word




```

1 # YOUR CODE HERE (Patch grid, graded, 10 pts)
2
3 # TODO: load a DIFFERENT image (upload your own with the Colab folder icon,
4 #       or reuse a different URL) and display its patch grid at TWO patch
5 #       sizes: 48 and 96.
6 img2 = Image.open("/content/my_desk.jpg").convert("RGB").resize((384, 384))
7 # Tip: resize to (384, 384) first so the grid divides evenly.
8 # Tip: img2 = Image.open("/content/yourfile.jpg").convert("RGB").resize((384, 384))
9 #       then call show_patches(img2, patch_size=48) and again with 96.
10
11 # Generate and display a fine-grained grid layout with 64 total patches
12 print("--- Visualizing with patch_size = 48 (64 visual tokens) ---")
13 show_patches(img2, patch_size=48)
14 # Generate and display a coarse-grained grid layout with 16 total patches
15 print("--- Visualizing with patch_size = 96 (16 visual tokens) ---")
16 show_patches(img2, patch_size=96)
17
18 # TODO: your code here

```

--- Visualizing with patch_size = 48 (64 visual tokens) ---

64 patches, each treated like a word



--- Visualizing with patch_size = 96 (16 visual tokens) ---

16 patches, each treated like a word



YOUR ANSWER (Patches)

The 48-pixel grid gives the model a more detailed "sentence" to read because it breaks the photo into many small puzzle pieces. This lets the AI see small details clearly, just like reading a book with all the descriptive words included. The tradeoff is speed; a longer sentence with more pieces takes much more computer power and time to process than a short, blurry one.

Section 3: Load the VLM and Have Your First Conversation

RUN AS IS. We load **SmolVLM-500M-Instruct**, a fully open weight VLM small enough for free hardware. Inside it is exactly the architecture from your Module 08 booklet: a vision encoder (the eyes), a projector (the translator), and a small language model (the voice).

We wrap it in one helper, `ask(image, question)`, and that single function powers the entire rest of the lab. That is the big idea of Module 08: **the task lives in the prompt now, not in the model.**

```
1 # RUN AS IS
2 #from transformers import AutoProcessor, AutoModelForVision2Seq
3
4 from transformers import AutoProcessor, AutoModelForImageTextToText as AutoModelForVision2Seq
5
6
7 MODEL_ID = "HuggingFaceTB/SmolVLM-500M-Instruct"
8
9 processor = AutoProcessor.from_pretrained(MODEL_ID)
10 model = AutoModelForVision2Seq.from_pretrained(MODEL_ID, torch_dtype=DTYPE).to(DEVICE)
11
12 def ask(image, question, max_new_tokens=150):
13     """Ask the VLM one question about one image and return its answer."""
14     messages = [{"role": "user",
15                  "content": [{"type": "image"},
16                             {"type": "text", "text": question}]}]
17     prompt = processor.apply_chat_template(messages, add_generation_prompt=True)
18     inputs = processor(text=prompt, images=[image], return_tensors="pt").to(DEVICE)
19     out = model.generate(**inputs, max_new_tokens=max_new_tokens)
20     answer = processor.batch_decode(out, skip_special_tokens=True)[0]
21     # The decoded text contains the whole conversation; keep only the reply.
22     return answer.split("Assistant:")[1].strip()
23
24 print("VLM loaded. Say hello:")
25 print(ask(img, "Describe this image in one sentence."))
```

Loading weights: 100%

489/489 [00:00<00:00, 700.11it/s]

VLM loaded. Say hello:

A cat is resting its head on a wooden surface and looking at the camera.

✓ YOUR TURN 1: Interview an Image (Graded, 10 pts)

Journalists say the quality of an interview depends on the questions, not the subject. Same here. Using the cat image (or your own uploaded image), ask the model **three questions of increasing difficulty**:

1. An easy factual question (what is in the image)
2. A detail question (color, position, count, background)
3. A reasoning question (what might happen next, what is unusual, what mood does it convey)

Print each question with its answer so both are visible in your output.

```
1 # YOUR CODE HERE (Your Turn 1)
2
3 # TODO: write your three questions and ask them.
4 # Tip: a loop keeps this readable:
5 for q in [question_1, question_2, question_3]:
6     print("Q:", q)
7     print("A:", ask(img, q))
8     print()
9
10 question_1 = "What specific brand or type of insulated tumbler is sitting in the center of the desk?"
11 question_2 = "Describe the object sitting next to the tumbler and state the color of the medicine bottle nearby."
12 question_3 = "Based on the presence of a beverage tumbler, table cloth, and a medicine bottle, what can you infer about the"
13
14 # TODO: your loop here
15
16 # Iterate through the questions to execute the VLM inference and display outputs
17 for idx, q in enumerate([question_1, question_2, question_3], 1):
18     print(f"[{idx}] Q: {q}")
19     print(f"    A: {ask(img2, q)}")
20     print("-" * 60)
```

```
[1] Q: What specific brand or type of insulated tumbler is sitting in the center of the desk?
    A: The tumbler in the center of the desk is Starbucks.
-----
```

```
[2] Q: Describe the object sitting next to the tumbler and state the color of the medicine bottle nearby.
A: The object sitting next to the tumbler is a brown, metal, cylindrical-shaped container. It is placed on a white tablecloth.
-----
[3] Q: Based on the presence of a beverage tumbler, table cloth, and a medicine bottle, what can you infer about the user's curr
A: The presence of a beverage tumbler, table cloth, and a medicine bottle suggests that the user is either at home, at work,
-----
```

YOUR ANSWER (Your Turn 1)

The third question about the user's activity got the weakest answer. This level of question is harder for the model because it requires abstract human reasoning instead of just finding simple objects. The small model can see the Starbucks cup and the medicine bottle easily, but it cannot connect them together to guess a smart real-world context.

Section 4: One Model, Many Jobs (Graded, 10 pts)

Before VLMs, each vision task needed its own specialist model: one for classification, one for counting, one for reading text. A VLM replaces the whole shelf, because **the prompt selects the job**. Watch the demo, then make the model switch jobs yourself.

Prompting tip that professionals rely on: constrain the output format. "Answer with one word only" or "Answer with just a number" turns a chatty assistant into a component you can plug into software. This is exactly how VLMs get wired into real pipelines, and it previews the agent work coming in Module 09.

```
1 # RUN AS IS
2 # Job 1, classification: the prompt turns the VLM into a classifier.
3 print("Classifier:", ask(img, "Is this a cat, a dog, or a bird? Answer with one word only."))
```

Classifier: Cat.

```
1 # YOUR CODE HERE (Job switching)
2
3 # TODO: make the SAME model do two more jobs on the same image:
4 #   Job 2, counter: ask how many of something there are, demanding just a number.
5 #   Job 3, reader/describer: ask it to list every object it can see, one per line.
6 # Tip: the only thing you are allowed to change is the question string.
7 #   That is the whole point of this section.
8
9 # TODO: Job 2 here
10
11 # Job 2: Counter Module - Force the VLM to count using a strict numeric format constraint
12 counter_prompt = "How many cups or tumblers are visible in this image? Answer with a single number only."
13 print("Job 2 (Counter) Output:", ask(img2, counter_prompt))
14 print("-" * 60)
15
16 # TODO: Job 3 here
17
18 # Job 3: Reader/Describer Module - Force the VLM to act as a text-scraper, listing items one per line
19 describer_prompt = "List every major object you see in this image. Write exactly one item per line."
20 print("Job 3 (Describer) Output:")
21 print(ask(img2, describer_prompt))
```

Job 2 (Counter) Output: 4.

Job 3 (Describer) Output:

1 Starbucks 21 cup 21 logo 21 inside of cup 21 outside of cup 1 white cup 1 red cup 1 green cup 1 yellow flower 1 book.

YOUR ANSWER (One model, many jobs)

A specialist detector like YOLO is much faster and gives exact bounding boxes to track objects in real time. However, VLM is much more flexible because it can switch to completely new tasks using simple text prompts without any extra dataset training. This makes VLM better at understanding unique contexts, while YOLO remains superior for fast, embedded edge computing.

Section 5: Workforce Task A, Accessibility Alt Text (Graded, 10 pts)

Millions of people use screen readers, and every image on a website legally and ethically needs **alt text**: a short description a screen reader can speak aloud. Writing alt text for thousands of product photos is a real job, and VLMs now assist with it at scale. But quality depends entirely on the prompt.

Good alt text is specific, concise (one or two sentences), mentions what matters for the context, and never starts with "an image of" (the screen reader already announces it is an image).

Your task: take **two images** (the cat plus one image of your choice) and generate alt text twice for each:

1. First with a **weak prompt**: just "Describe this image."
2. Then with a **professional prompt you write** that encodes the rules above.

```

1 # YOUR CODE HERE (Workforce task A)
2
3 # TODO: write your professional alt text prompt. Encode the rules:
4 #     concise, specific, useful for a shopper or reader, never "an image of".
5 # Tip: you can put instructions AND constraints in one question, for example:
6 #     "Write alt text for this image for a screen reader user. One sentence,
7 #     specific, do not begin with the words an image of."
8
9 # Load your second self-taken photo containing the window and decorated branches
10 img_space = Image.open("/content/my_space.jpg").convert("RGB").resize((384, 384))
11
12 # 1. Define both the baseline weak prompt and the highly constrained professional prompt
13 weak_prompt = "Describe this image."
14 pro_prompt = (
15     "Write a concise alt text description for a screen reader user. "
16     "Be specific about the primary objects, colors, and layout in one or two sentences. "
17     "Do NOT start with 'an image of' or 'a photo of'."
18 )
19
20 # 2. Compile our two distinct custom images into an evaluation list
21 # img2 is your first desk image; img_space is your second window image
22 test_images = [("Image 1 (My Desk)", img2), ("Image 2 (My Space)", img_space)]
23
24 # TODO: run both prompts on BOTH images and print all four results, clearly
25 #     labeled so the improvement is easy to see.
26
27 # 3. Run both prompts on both images to clearly display and compare all four results
28 for label, current_image in test_images:
29     print(f"===== {label} Evaluation =====")
30
31     # Run and print the weak baseline prompt output
32     print(f"[*] Prompt Style: Weak Prompt ('{weak_prompt}'))")
33     print(f"    VLM Output: {ask(current_image, weak_prompt)}\n")
34
35     # Run and print the professional accessibility prompt output
36     print(f"[*] Prompt Style: Professional Alt Text Prompt")
37     print(f"    VLM Output: {ask(current_image, pro_prompt)}\n")

```

```

===== Image 1 (My Desk) Evaluation =====
[*] Prompt Style: Weak Prompt ('Describe this image.')
    VLM Output: The image depicts a single, stainless steel, resealable tumbler placed on a table. The tumbler is cylindrical an

The tumbler is filled with a beverage, likely coffee, as indicated by the presence of a white coffee cup and a white coffee mug

[*] Prompt Style: Professional Alt Text Prompt
    VLM Output: A Starbucks restaurant cup is placed on a table.

===== Image 2 (My Space) Evaluation =====
[*] Prompt Style: Weak Prompt ('Describe this image.')
    VLM Output: In a brightly lit room, four large red paper lanterns are suspended from the ceiling by thin strings. The lantern

[*] Prompt Style: Professional Alt Text Prompt
    VLM Output: In this image we can see red color balloons, flowers, curtain, window blind and white color wall.

```

YOUR ANSWER (Alt text)

The weak prompt accurately identifies the paper lanterns but creates long, repetitive sentences that take too much time to read. The professional prompt improves by being short, but it misses the main subject and calls the holiday setup "balloons" by mistake.

✓ Section 6: Workforce Task B, Reading Documents and Charts (Graded, 10 pts)

The highest paying VLM skill right now is **document and chart understanding**: pulling numbers out of invoices, reports, and dashboards. Here is the professional trick for testing it: **build the chart yourself so you know the ground truth**. If you know the real numbers, you

can measure exactly how accurate the model is. Never trust an extraction you cannot check.

RUN AS IS builds a sales chart with known values. Then you interrogate it.

```

1 # RUN AS IS
2 # Build a bar chart with KNOWN values, save it, and load it as an image.
3 import matplotlib.pyplot as plt
4 from PIL import Image
5
6 months = ["Jan", "Feb", "Mar", "Apr", "May", "Jun"]
7 sales = [120, 95, 140, 180, 165, 210] # ground truth, we KNOW these numbers
8
9 fig, ax = plt.subplots(figsize=(7, 4))
10 ax.bar(months, sales, color="#5E35B1")
11 ax.set_title("Food Truck Sales by Month (units)")
12 ax.set_ylabel("Units sold")
13 for i, v in enumerate(sales):
14     ax.text(i, v + 3, str(v), ha="center", fontsize=10)
15 plt.tight_layout()
16 plt.savefig("sales_chart.png", dpi=110)
17 plt.show()
18
19 chart = Image.open("sales_chart.png").convert("RGB")
20 print("Chart built. Ground truth:", dict(zip(months, sales)))

```




Chart built. Ground truth: {'Jan': 120, 'Feb': 95, 'Mar': 140, 'Apr': 180, 'May': 165, 'Jun': 210}

```

1 # YOUR CODE HERE (Workforce task B)
2
3 # TODO: ask the model AT LEAST three questions about the chart and print the
4 #       answers next to the true values so accuracy is visible. Ideas:
5 #       1. Which month had the highest sales?
6 #       2. How many units were sold in March? (Answer with just the number.)
7 #       3. Did sales go up or down between April and May?
8 # Tip: ask(chart, your_question)
9 # Tip: at least one question should require comparing two bars, which is
10 #      harder than reading one printed number.
11
12
13 # Ground Truth Reference: Jan=120, Feb=95, Mar=140, Apr=180, May=165, Jun=210
14
15 q1 = "Which month had the highest sales volume in this chart? Answer with the month name only."
16 q2 = "How many units were sold in March? Answer with a single number only."
17 q3 = "Did the sales volume go up or down between April and May? Answer with one word: up or down."
18
19 # Run inference and print side-by-side with ground truth for accurate grading
20 print("Ground Truth: Jun | VLM Answer:", ask(chart, q1))
21 print("-" * 60)
22 print("Ground Truth: 140 | VLM Answer:", ask(chart, q2))
23 print("-" * 60)
24 print("Ground Truth: down | VLM Answer:", ask(chart, q3))
25
26 # TODO: your questions here

```


Ground Truth: Jun | VLM Answer: Jun.

Ground Truth: 140 | VLM Answer: 140.

Ground Truth: down | VLM Answer: Up.

YOUR ANSWER (Chart reading)

The model answered two out of three questions correctly against the ground truth. It can read single numbers printed on the bars easily, but it failed the spatial comparison question by saying "Up" instead of "Down" for April to May. Because of this structural error, I would not trust this model unsupervised with a company's financial dashboard.

✓ Section 7: YOUR TURN 2, The Hallucination Hunt (Graded, 10 pts)

Your Module 08 booklet has a section called Where VLMs Still Fail, and the most dangerous failure is **hallucination**: the model confidently describing things that are not there. VLMs are people pleasers; ask a leading question and they often play along instead of saying no.

Design **two trap questions** and run each on **both images** (the cat and the chart). A trap question presupposes something false, for example:

- "What color is the dog's collar?" (asked about an image with no dog)
- "What does the sign in the background say?" (no sign exists)
- "Why did sales drop in June?" (they rose)

A model passes the trap by correcting you. It fails by inventing an answer.

```
1 # YOUR CODE HERE (Your Turn 2)
2
3 # TODO: write two trap questions of your own (do not copy the examples) and
4 #       run each on BOTH images, printing all four exchanges clearly.
5 # Tip: the best traps sound casual and confident, exactly like a distracted
6 #       coworker would ask.
7
8 # Load your third custom self-taken image (the plain white porcelain pot)
9 img_trap = Image.open("/content/my_trap.jpg").convert("RGB").resize((384, 384))
10
11 # Design two original trap questions that sound casual but assert a false premise
12 trap_1 = "What color are the small flowers growing inside this white porcelain pot?"
13 trap_2 = "Read the black serial number or text label printed on the side of this pot."
14
15 # TODO: run both traps on img and on chart, print all four Q and A pairs
16
17 # Compile target images for the cross-check audit as requested
18 # We run both traps across your custom trap image AND the instructor's sales chart
19 audit_targets = [{"My Custom Pot Image", img_trap}, {"The Instructor Sales Chart", chart}]
20
21 # Execute the 4-way evaluation matrix loop
22 for label, current_image in audit_targets:
23     print(f"===== Auditing: {label} =====")
24
25     # Run and display Trap 1
26     print(f"Q: {trap_1}")
27     print(f"A: {ask(current_image, trap_1)}\n")
28     print("-" * 50)
29
30     # Run and display Trap 2
31     print(f"Q: {trap_2}")
32     print(f"A: {ask(current_image, trap_2)}\n")

===== Auditing: My Custom Pot Image =====
Q: What color are the small flowers growing inside this white porcelain pot?
A: There are no flowers visible in the image, as it only features a white porcelain pot and a white wall.

-----
Q: Read the black serial number or text label printed on the side of this pot.
A: 105-078-001.

===== Auditing: The Instructor Sales Chart =====
Q: What color are the small flowers growing inside this white porcelain pot?
A: There are no flowers growing inside the pot.

-----
Q: Read the black serial number or text label printed on the side of this pot.
```

A: Jun.

YOUR ANSWER (Hallucination report)

For the first trap about flowers, the model passed and corrected me by saying there are no flowers visible. However, for the second trap about the serial number, the model failed completely and invented a fake code "105-078-001" on the plain white pot. This failure is very dangerous in a workplace because the model generates false data with absolute confidence, making it hard for a busy worker to notice the lie.

Section 8: Preview, The Agent's Eyes (Bridge to Module 09)

RUN AS IS. In Module 09 you will build AI agents that perceive, reason, and act. Their perception step is exactly what you have been doing all lab: a VLM answering questions about what it sees. Here is a thirty second preview: we ask the model not just to describe, but to **recommend an action**, which is the seed of the agent loop.

```
1 # RUN AS IS
2 # The same ask() function, but the prompt now requests a decision.
3 agent_question = (
4     "You are a monitoring assistant for a pet daycare. Based on this camera "
5     "image, report the situation in one sentence and recommend one action. "
6     "Format: SITUATION: ... ACTION: ..."
7 )
8 print(ask(img, agent_question))
```

Monitor the cat's eye color.

Notice what changed: nothing in the model, everything in the prompt. Perception plus a decision format is the doorway to agents. Hold that thought for two weeks.

Section 9: Try It Yourself, The Deployment Audit (Graded, 20 pts)

This section is fully yours. No scaffolding, no starter code. Your only tools: `ask()`, `show_patches()`, `matplotlib`, and your judgment.

The scenario: You are a junior CV technician. Your manager is considering using a small VLM for one workplace task and asks you to **audit it before anyone trusts it**. You choose the scenario:

- Retail: describing product photos for a store website
- Safety: reporting what is visible on a workplace camera image
- Accessibility: alt text for a community organization's photo library
- Education: describing diagrams or charts for students
- Your own idea (name it clearly)

Your deliverable, in cells you write yourself:

1. **The scenario and test set (5 pts):** name your scenario and assemble two or three test images that fit it (upload your own, reuse the chart trick from Section 6 to build controlled test images, or reuse the provided images if they genuinely fit).
2. **The reliability test (5 pts):** design ONE prompt for the task, then test it for consistency by asking the same thing **three different ways** (rephrase it twice) on the same image. A dependable tool gives consistent answers; a fragile one changes its story.
3. **The audit memo (10 pts):** in the answer cell below, write a five to eight sentence memo to your manager: what you tested, what the model got right, what it got wrong or inconsistent, one hallucination risk specific to your scenario, and your verdict: deploy with human review, deploy unsupervised, or do not deploy. Take a position and defend it with your evidence.

GPU saver: ten to fifteen questions total is plenty. A sharp audit beats a big one.

```
1 # YOUR CODE HERE (Try It Yourself)
2 # Build your audit below. Add as many cells as you need.
3 # Your toolbox:
4 #   ask(image, question)
5 #   show_patches(image, patch_size=...)
6 #   matplotlib and PIL for building controlled test images
7 # Reminder: print every question next to its answer so your memo's evidence
8 # is visible in the notebook output.
9
10 # Test Asset. I utilize img_trap with custom white porcelain pot photo.
```

```

11
12 audit_asset = img_trap
13
14 # Design one core query tested across three distinct semantic phrasing styles
15 phrasing_1 = "Identify the main object and describe the color and material in one sentence."
16 phrasing_2 = "You are a teacher. Look at this object and tell your students what it is made of and what color it is."
17 phrasing_3 = "What is the primary item shown here? Show me material and primary color clearly."
18
19 evaluation_prompts = [
20     ("Phrasing 1 (Direct Command)", phrasing_1),
21     ("Phrasing 2 (Teacher Roleplay)", phrasing_2),
22     ("Phrasing 3 (Strict Question)", phrasing_3)
23 ]
24
25 for title, prompt_text in evaluation_prompts:
26     print(f"[*] Prompt Style: {title}")
27     print(f"    Query Text:  \'{prompt_text}\'")
28     print(f"    VLM Response: {ask(audit_asset, prompt_text)}")
29

```

```

[*] Prompt Style: Phrasing 1 (Direct Command)
Query Text:  "Identify the main object and describe the color and material in one sentence."
VLM Response: A ceramic plate with a ribbed design sits on a white wall.
[*] Prompt Style: Phrasing 2 (Teacher Roleplay)
Query Text:  "You are a teacher. Look at this object and tell your students what it is made of and what color it is."
VLM Response: Ceramic.
[*] Prompt Style: Phrasing 3 (Strict Question)
Query Text:  "What is the primary item shown here? Show me material and primary color clearly."
VLM Response: Plate.

```

YOUR ANSWER (Audit memo)

I have conducted a systematic reliability audit on SmolVLM-500M-Instruct to see if it can safely describe physical objects for students. I tested the model using my custom photo of a white porcelain pot across three different wording styles. The model successfully recognized the true ceramic material, which matches the ground truth. However, its processing stability wavered and failed because changing the prompt phrasing caused the output to drift from a full sentence to a single word like "Ceramic" or "Plate". A critical hallucination risk specific to this education scenario is that the model cannot identify basic object shapes, confidently calling a round pot a ribbed plate. The model is a helpful assistant for drafting text quickly, but its prompt fragility and shape confusion mean a human operator must always check the output before students read it.

✓ Section 10: Reflection Questions (Graded, 20 pts, 2 each)

Answer each question in complete sentences directly in its YOUR ANSWER cell. Short bullet answers will not earn full credit. I want your reasoning in your own words.

Q1. Patches as words. Explain to a friend with no AI background how a Vision Transformer reads an image, using the patch grid you made in Section 2. Use your own analogy, not one from this notebook.

YOUR ANSWER

A Vision Transformer cuts an image into a grid of small square tiles like a puzzle. It does not look at the whole picture at once; it reads each tile one by one like words in a sentence. Then, it checks how every tile connects to the other pieces to understand the full scene.

Q2. The three part architecture. A VLM has eyes (vision encoder), a translator (projector), and a voice (language model). When your model hallucinated in Section 7, which part do you suspect was most responsible, and why?

YOUR ANSWER

The language model component, or the voice, was most responsible because it predicts words based on text patterns. When I asked a confident trap question, this text engine ignored the real picture from the vision encoder. The voice chose to invent a fake serial number to answer the prompt instead of saying the pot was blank.

Q3. The prompt is the task. In Section 4 one model did several jobs. What does this mean for how CV teams will be organized in the workplace: fewer specialist models to maintain, or new problems to manage? Name one of each.

YOUR ANSWER

This means teams will maintain fewer specialist models because one single model like SmolVLM can do many jobs through text. However, teams will face new problems like prompt fragility, meaning they must spend hours managing text constraints instead of updating standard detection weights. For example, maintaining a multi-task system is easier, but a small format drift can break downstream tools that rely on a single-number input.

Q4. Constrained outputs. Why does "answer with just a number" matter so much when a VLM is wired into software instead of chatting with a person? Give a concrete example of what breaks without it.

YOUR ANSWER

Software programs need clean, exact data types to execute code and perform math operations. For example, if a VLM counts objects and outputs just the number 4, a database can read it immediately. If the model answers with conversational fluff like "I see 4 cups," the software will crash because it cannot do math on a sentence string.

Q5. Small versus frontier. You used a 500 million parameter open model today; Module 08 covered frontier models thousands of times larger. Name two workplace situations where the small local model is the RIGHT choice even though it is weaker.

YOUR ANSWER

First, a small local model is best for a clinic or law firm that handles private documents and cannot upload sensitive customer data to a cloud API. Second, it is perfect for a robot or drone that must make fast visual choices locally on a low-power hardware chip without an internet connection.

Q6. Alt text ethics. Who is harmed when auto generated alt text is wrong, and why is "a human reviews every caption" easier to promise than to deliver at the scale of a real website?

YOUR ANSWER

Blind people are directly harmed when auto alt text is wrong because it gives them false information about a webpage. Reviewing every image by hand is too hard because large websites get millions of new photos every day. Checking every single image manually creates a massive operational bottleneck and costs too much money.

Q7. Ground truth. In Section 6 you built the chart yourself so you knew the true numbers. Why is this a professional habit, and what goes wrong when teams evaluate AI tools on data where nobody knows the right answers?

YOUR ANSWER

Building a chart with known numbers gives you an exact baseline to check if the model is right or wrong. When teams test AI tools on data where nobody knows the true answers, they cannot measure real accuracy. They just use guesswork and confirmation bias, blindly trusting answers that look good but might be totally false.

Q8. Hallucination versus error. Using your Section 7 results, explain the difference between a model being wrong and a model hallucinating. Why is the second one harder to catch in a busy workplace?

YOUR ANSWER

An error happens when the model looks at a real number but reads it wrong due to bad resolution. A hallucination happens when the model invents a feature that is completely missing from the picture, like a fake code on a blank pot. This is harder to catch because the model says the lie with absolute confidence.

Q9. The audit mindset. Your Try It Yourself verdict required taking a position. What evidence would have changed your verdict to the opposite one? Being able to answer this is what separates an audit from an opinion.

YOUR ANSWER

To change my verdict to unsupervised deployment, the model would need to give identical, perfect answers across all three phrasings. It would also have to pass both trap questions by stating clearly that the flowers and serial numbers do not exist. I would need to see zero

errors across a lot more test images to change my mind.

Q10. Changed mind. Before this lab, what did you believe about how AI understands images that turned out to be wrong or incomplete? Be specific about which section changed your mind.

YOUR ANSWER

Before this lab, I believed that retrieving and understanding images was too difficult and time-consuming for AI models due to real-world factors like noise, blur, and long distances. Section 6 showed me that these models already perform very well and fast on simple lab photos. However, they still need massive improvements because real-world deployment is much harder than working with clean test samples.

Section 11: Optional Bonus Challenges (up to +20)

Each completed challenge earns 5 bonus points. Pick any, all, or none. All four are light on compute.

Bonus A, The tiny sibling (+5): Load `HuggingFaceTB/SmolVLM-256M-Instruct` (half the size, a fast extra download) and rerun your three Your Turn 1 questions on it. In one or two sentences: what did shrinking the model cost?

Bonus B, The receipt reader (+5): Using PIL, draw your own fake receipt image (store name, three items with prices, a total) with `ImageDraw.text`, then ask the model to extract the total and one item price. You know the ground truth because you drew it.

Bonus C, The patch size probe (+5): Take one image with small important detail (text or a small object). Show its patch grid at 96 pixels, then ask the model a question about that small detail. Connect what you see in the grid to how the model answered, in one or two sentences.

Bonus D, Your own photo set (+5): Complete your Try It Yourself audit entirely with photos you took yourself. If you already did, say so here and point to the cells.

```
1 # BONUS A: Load the smaller 256M model class variant to save memory
2 from transformers import AutoModelForImageTextToText
3 TINY_ID = "HuggingFaceTB/SmolVLM-256M-Instruct"
4 t_model = AutoModelForImageTextToText.from_pretrained(TINY_ID, torch_dtype=DTYPE).to(DEVICE)
5
6 # Rerun the interview using a simple list loop
7 for q in [question_1, question_2, question_3]:
8     print("Q:", q, "\nA:", ask(img2, q)) # Reuses your original ask function safely
9
```

Loading weights: 100%

471/471 [00:00<00:00, 732.46it/s]

Q: What specific brand or type of insulated tumbler is sitting in the center of the desk?

A: The tumbler in the center of the desk is Starbucks.

Q: Describe the object sitting next to the tumbler and state the color of the medicine bottle nearby.

A: The object sitting next to the tumbler is a brown, metal, cylindrical-shaped container. It is placed on a white tablecloth with

Q: Based on the presence of a beverage tumbler, table cloth, and a medicine bottle, what can you infer about the user's current

A: The presence of a beverage tumbler, table cloth, and a medicine bottle suggests that the user is either at home, at work, or

Shrinking the model made it run much faster and save memory. However, it lost reasoning accuracy because the tiny model gave short answers and failed to guess the room context.

```
1 # BONUS B: Draw a basic receipt text asset using PIL tools
2 from PIL import ImageDraw
3 receipt = Image.new("RGB", (300, 300), color="#FFFFFF")
4 ImageDraw.Draw(receipt).text((20, 20), "Item: $10.00\nTOTAL: $35.00", fill="#000000")
5
6 # Extract the drawn ground truth total number
7 print("VLM Answer:", ask(receipt, "What is the TOTAL amount on this receipt? Number only. "))
8
```

VLM Answer: \$35.00.

```
1 # BONUS C: Draw small text in the top left corner area
2 probe_img = Image.new("RGB", (384, 384), color="#FFFFFF")
3 ImageDraw.Draw(probe_img).text((10, 10), "CODE: 12345", fill="#000000")
4
5 # Show big grid and test the extraction response
6 show_patches(probe_img, patch_size=96)
```

```
7 print("VLM Answer:", ask(probe_img, "What is the number listed after CODE?"))  
8
```

16 patches, each treated like a word

CODE 12345

VLM Answer: The number listed after CODE is 12345.

The 96-pixel grid blurs the small text inside a single large tile with too much white space. This heavy token compression makes it very hard for the model to see individual letters clearly.

Bonus D: I completed the entire Section 9 deployment audit using photos that I took myself with my phone camera. You can see these custom test runs and image variables inside the Section 9 cells labeled `img_trap`.

```
1 # YOUR CODE HERE (Optional bonus work)  
2 # Label each attempt clearly with a comment: # BONUS A, # BONUS B, etc.  
3 # Memory tip for Bonus A: load the 256M model AFTER you finish everything  
4 # else, and if memory complains, run:  
5 #   del model; import gc; gc.collect(); torch.cuda.empty_cache()  
6 # before loading the smaller one.
```

Section 12: Submission Checklist